

# skaters: Model First, Conform Last — Composable Online Distributional Time-Series Forecasting in Python and JavaScript

Peter Cotton

Microprediction

[peter.cotton@microprediction.com](mailto:peter.cotton@microprediction.com)

<https://www.microprediction.org>

June 18, 2026

**Abstract.** We present **skaters**, a zero-dependency library for online univariate time-series forecasting in which *every prediction is a full probability distribution*, not a point. Models are built by composition: invertible transforms chain together, ensembles nest, and a single distributional *leaf* sits at the bottom of every chain. The leaf is a pluggable proper-scoring-rule optimiser, which lets us separate two concerns that the forecasting and conformal-prediction literatures usually conflate — *modelling* the conditional mean (judged by likelihood) and *calibrating* the predictive tail (optionally judged by a downstream score such as the continuous ranked probability score, CRPS). We call the resulting recipe **model first, conform last**, and show that it matches a dedicated CRPS specialist on CRPS while *improving* held-out log-likelihood. In a fair rolling one-step-ahead comparison on 500 FRED change-series against *eight* distributional baselines — **AutoARIMA**, **AutoETS**, **statsmodels SARIMAX** and exponential smoothing, a **GARCH(1,1)** model with Student-*t* innovations, a **NeuralForecast** model with a Student-*t* head, and **AutoARIMA** paired with split-conformal and adaptive-conformal residuals — the general forecaster **wins the per-series log-likelihood race against all eight**. The decisive test is the **GARCH-*t*** model, the classical heavy-tail / volatility-clustering state of the art: our forecaster still edges it on the continuous subset (mean continuous log-likelihood 2.86 versus 2.72), confirming the heavy-tail advantage *against the heavy-tail specialist*. On CRPS the picture is mixed, as expected: we beat the smoothing, **GARCH** and neural baselines, roughly tie the **ARIMA** family, and lose to the CRPS-optimised conformal variants. We further contribute a *mean-preserving lattice projection* for series that revisit exact values, an online *coordinate* (Yeo–Johnson) search, and a committed *martingale* model whose volatility clock is a time-changed Brownian motion. The whole library is implemented twice — pure Python and zero-dependency JavaScript — and verified identical to  $10^{-6}$  on roughly 90,000 probe values, so the same models run server-side or in the browser via **Pyodide**.

**Keywords:** probabilistic forecasting, time series, proper scoring rules, conformal prediction, online learning, Python, JavaScript, **skaters**.

## 1 Introduction

Most deployed forecasters return a point and, perhaps, an interval bolted on afterwards. Yet the quantity a decision-maker actually needs is the *predictive distribution*: the mean for an estimate, the spread for risk, the tails for position sizing, and a density for likelihood-based evaluation and combination. Two strands of recent work take distributions seriously but each gives something up.

Classical state-space and exponential-smoothing models [Hyndman et al., 2002, Hyndman and Khandakar, 2008] are distributional but rarely *online and composable*: adding

a transform usually means re-deriving the filter. Conditional-heteroscedasticity models [Engle, 1982, Bollerslev, 1986] capture the heavy tails and volatility clustering of financial data, but commit to a single parametric form and are refit in batch. At the other extreme, **conformal prediction** [Vovk et al., 2005, 2019] is distribution-free and online, but a conformal predictive system outputs a *cumulative distribution function*, not a density. It is calibrated for coverage and CRPS but **cannot be scored on log-likelihood at all**, and an empirical conformal CDF assigns zero density — and hence  $-\infty$  log-likelihood — to any value outside the observed residual range. Modern neural and foundation models [Salinas et al., 2020, Ansari et al., 2024, Das et al., 2024, Woo et al., 2024] are powerful distributional learners, but are built for global cross-series training on hardware accelerators, not for a single univariate stream running in a browser.

**skaters** takes a third route. Every node — transform, ensemble, or leaf — is a function  $(y, \text{state}) \rightarrow ([\text{Dist}], \text{state})$  that consumes one observation and emits a list of **Dist** objects, one per forecast horizon, where a **Dist** is a weighted Gaussian mixture. Because the type is closed under composition, an entire model is a chain of invertible transforms with one distributional leaf at the bottom, and ensembles are simply skaters that combine other skaters. This single abstraction yields three contributions developed in turn:

1. **A pluggable objective** (Sections 3–4). The leaf fits its mixture weights by a proper scoring rule. Point it at log-likelihood and it models honestly; point its *terminal* copy at CRPS while keeping a likelihood-weighted trunk, and you get **model first, conform last** — competitive with a CRPS specialist on CRPS *and* better on likelihood.
2. **Distributional structure that survives composition** (Sections 5–7). A *mean-preserving lattice projection* places atoms on the exact values a series revisits without moving the ensemble mean and while vanishing on continuous data; an online *co-ordinate* search learns whether a series is simple in a log, root, or linear coordinate; and a committed *martingale* model learns only a volatility clock.
3. **An honest benchmark against eight distributional baselines** (Section 8), including the heavy-tail (**GARCH-*t***) and neural (**NeuralForecast**) state of the art, scored on held-out log-likelihood and CRPS in a single, identical **Dist**-based protocol.

The unifying premise is the No-Free-Lunch theorem [Wolpert and Macready, 1997] taken seriously: every “method” is a prior, and averaged over all series none dominates. The right response is not to pick one but to make the priors explicit, weight them online, and bound the cost of a wrong prior by how fast the weighting abandons it.

## 2 The **Dist** abstraction and composable transforms

A **Dist** is a weighted mixture  $\sum_i w_i \mathcal{N}(\mu_i, \sigma_i^2)$  with  $\sigma_i > 0$  and  $\sum_i w_i = 1$ . It exposes **mean**, **std**, **pdf**, **cdf**, **logpdf**, **quantile**, and a closed-form **crps**, together with the affine maps (**shift**, **scale**, **affine**) and a moment-matching **prune** needed to propagate it through transform inverses and to bound component growth in ensembles. The Gaussian mixture is a deliberate choice: it is closed under affine maps and mixing, admits a closed-form CRPS through the elementary identity  $\mathbb{E} |\mathcal{N}(m, s^2)| = 2s \phi(m/s) + m(2\Phi(m/s) - 1)$ , and approximates heavy tails arbitrarily well as a scale mixture.

A *transform* is an online bijection given by a **forward** (scalar in, scalar out) and an **inverse\_k** that maps  $k$  **Dists** in the transformed space back to the original. Differencing,

exponential smoothing, drift, Holt’s linear trend, fractional differencing, autoregression (via online recursive least squares), **GARCH**-style scaling, seasonal differencing, power and Yeo–Johnson coordinates [Box and Cox, 1964, Yeo and Johnson, 2000], and standardisation are all transforms. *Conjugation* composes a transform with an inner skater,

$$y \xrightarrow{T} y' \rightarrow \text{inner} \rightarrow D' \xrightarrow{T^{-1}} D,$$

and because every step is online and the inverse propagates the full mixture, multi-step uncertainty accumulates correctly — for example variance grows as  $\sum_h \sigma_h^2$  under differencing.

**Numerical note.** Mixture variance is computed in centred form,  $\text{Var} = \sum_i w_i (\sigma_i^2 + (\mu_i - \bar{\mu})^2)$ , rather than  $\mathbb{E}[X^2] - \mathbb{E}[X]^2$ . The latter catastrophically cancels when a tight component sits at a large mean, silently producing  $\sigma = 0$  and an invalid predictive. Running moments use Welford’s recurrence [Welford, 1962].

### 3 The terminal-leaf ensemble: model first

Bayesian model averaging over a heterogeneous pool [Hoeting et al., 1999, Raftery et al., 2005] combines full predictive distributions, which preserves the combined mean and variance but *washes out higher moments*: a heavy-tailed leaf used inside the pool collapses toward a Gaussian shape at the output, because a mixture of differently located Gaussians is dominated by the spread of their means. We therefore use the candidate models only as **mean forecasters**, weight them online by predictive likelihood (with a learning-rate shrinkage and a per-depth complexity penalty, in the spirit of gradient boosting), combine their means, and model the distribution of the *combined residual* with a **single terminal leaf**:

$$y \rightarrow \hat{\mu} = \sum_i w_i \hat{\mu}_i, \quad r = y - \hat{\mu} \rightarrow \text{leaf} \rightarrow D, \quad \hat{D} = D \text{ shifted by } \hat{\mu}.$$

Because exactly one leaf reaches the output, its shape — heavy tails, learned online — survives undiluted. This is the *model first* half of the recipe: a likelihood-weighted trunk whose only job is to get the conditional mean right. The candidate pool spans exponential smoothing, differencing, drift, Holt’s linear trend, autoregression, grouped long-lag autoregression, fractional differencing, a seasonal block, a Yeo–Johnson coordinate grid (Section 6), and a fast/slow two-systems block.

### 4 Metric-agnostic leaves: conform last

The leaf is itself a fixed Gaussian *scale mixture* — components  $\mathcal{N}(0, c_k \hat{\sigma})$  over a fixed log-spaced dictionary  $\{c_k\}$ , with  $\hat{\sigma}$  an exponentially weighted scale — whose only learned quantities are the mixture weights  $\{w_k\}$  on the simplex. A Gaussian scale mixture is a Student-*t* in the limit, so the leaf approximates heavy tails by construction while remaining a plain Dist. The weights are fit online by a **proper scoring rule** [Gneiting and Raftery, 2007]:

- **likelihood** — recency-weighted expectation–maximisation on the component responsibilities;

- **crps** — exponentiated-gradient descent on the simplex minimising the closed-form mixture CRPS [Matheson and Winkler, 1976]; the gradient is exact, since the mixture CRPS is a finite sum of  $\mathbb{E}|\mathcal{N}(m, s^2)|$  terms.

Placing the CRPS leaf at the *terminal* of a likelihood-weighted trunk gives **model first, conform last**. The trunk’s job is to model, and the honest objective for modelling is likelihood — a proper, tail-sensitive score that is, up to a constant, the Kelly log-growth criterion [Kelly, 1956]; corrupting it to chase CRPS would repeat the conformal mistake of discarding the density. The tail’s job is to be shaped, and that is the only place a downstream metric (CRPS) should get a vote. Section 9 shows this configuration dominates a pure-likelihood policy on CRPS *and* on likelihood: a CRPS-fit leaf is more outlier-robust, so it even generalises slightly better in log-likelihood. The objective is a knob, not a fixed cost.

## 5 The lattice projection

Administrative and grid-quoted series — policy rates, posted prices, pegged exchange rates — revisit a small set of exact values, and a continuous predictive density cannot place mass there. We add a *projection*: a recency-weighted frequency table of exact values, with near-Dirac atoms on the values revisited above a noise floor, each carrying its observed frequency as probability. The atom is **mean-preserving**: the continuous part is recentred so the atoms add *mass* without moving the ensemble mean.

Two properties make this a safe default. First, on continuous data nothing is revisited, no atom fires, and the projection vanishes identically — it is free when unused. Second, unlike a last-value spike it captures values that recur often but never consecutively (a rate that repeatedly returns to a round number). The pull of the level (its tendency to persist) is left to the trunk, where the random-walk candidate earns its weight by likelihood, so the projection remains a pure statement about mass, not mean.

## 6 Coordinate learning

Whether a series is simple in a linear, log/multiplicative, or root coordinate is itself learnable. We add a coarse grid of **Yeo–Johnson**  $\lambda$  candidates — the signed Box–Cox family [Box and Cox, 1964, Yeo and Johnson, 2000], defined on all of  $\mathbb{R}$  — to the pool, so the ensemble learns the coordinate online rather than committing up front. This is No-Free-Lunch-safe: a wrong coordinate is simply down-weighted by likelihood. The same mechanism expresses a non-negativity prior at  $\lambda = 0$  (the log coordinate), though we note that because the **Dist** is a Gaussian mixture the inverse uses the delta-method linearisation, so non-negativity is approximate rather than exact.

## 7 A martingale with a learned volatility clock

For near-martingale *levels* we provide a committed model, **doob**: the mean is pinned to the last value (a driftless martingale) and only the volatility clock is learned, as a Bayesian average over martingale predictives that differ in their volatility model (constant, **GARCH**, slowly varying, heavy-tailed). Because every candidate shares the same mean, plain averaging blends the clocks *without* washing out kurtosis — the committed mean is exactly what makes model averaging the right tool here, in contrast to Section 3. By the Dambis–Dubins–Schwarz theorem [Dambis, 1965, Dubins and Schwarz, 1965, Doob, 1953] a continuous martingale is a time-changed Brownian motion, so the model is literally

“Brownian motion on a stochastic clock”. It beats the general forecaster on near-martingale levels by committing the mean and spending its capacity on the clock; on genuinely mean-reverting series the martingale prior is wrong and it gives ground — a deliberately sharp instrument, fed *levels* rather than pre-differenced changes.

## 8 Experiments

**Universe.** To avoid a hand-picked series list we draw from the Federal Reserve Economic Data [Federal Reserve Bank of St. Louis], taking cached daily series and auto-transforming each to its one-step *changes* (log-difference for strictly positive levels, else first difference) — the stationary-ish innovation stream where heavy tails and regime shifts live, and which keeps log-likelihood comparable across series of different scale. We keep the last 1000 changes per series and score the last 300 as a rolling one-step-ahead test window after a burn-in.

### 8.1 Against eight distributional baselines

On 500 such series we compare the general forecaster **laplace** against every distributional one-step baseline we could assemble: **statsforecast**’s [Garza et al., 2022] **AutoARIMA** and **AutoETS**; **statsmodels** **SARIMAX**(1,0,1) and simple exponential smoothing, which emit *exact* closed-form Gaussian predictives; a constant-mean **GARCH**(1,1) model with Student-*t* innovations [Bollerslev, 1986, 1987] from the **arch** package, the classical heavy-tail / volatility-clustering state of the art; a **NeuralForecast** [Olivares et al., 2022] multilayer perceptron with a Student-*t* distribution head; and **AutoARIMA** paired with conformal residual quantiles [Vovk et al., 2019] in both a split-conformal and an adaptive-conformal (ACI, Gibbs and Candès, 2021) variant.

The protocol is a *fair rolling one-step-ahead*: each baseline is fit on an expanding or rolling window with periodic refit, and every method — ours and theirs — is turned into the *same* **Dist** and scored on held-out log-likelihood *and* CRPS over the same test window. The parametric Student-*t* predictives (**GARCH**, **NeuralForecast**) are represented as finite Gaussian scale mixtures, matching the analytic *t* density to  $\sim 10^{-3}$ . We report **per-series** win-rates (the fraction of series on which **laplace** scores better), and additionally restrict to the 309 *continuous* series — those with under 5% exactly-repeating one-step changes — since on repeating/grid series the lattice projection (Section 5) gives a large but metric-specific advantage that would otherwise dominate the aggregate. Table 1 reports the result.

**The likelihood race.** **laplace** wins the per-series log-likelihood race against *all eight* baselines, on the full universe and on the continuous subset, including the two exact-Gaussian references (**SARIMAX**, statsmodels ETS) for which there is no band-reconstruction approximation to blame.

**GARCH-*t* is the decisive test.** The **GARCH**(1,1)-*t* model is purpose-built for the heavy tails and volatility clustering of financial change-series, so it is *supposed* to be the hardest opponent on log-likelihood — and it is. It is the tightest race (67% of continuous series), and its mean continuous log-likelihood, 2.72, is closest to ours, 2.855 — a genuine but *narrow* edge of about 0.14 nats per observation. That the advantage survives *against the heavy-tail specialist* is the result that matters: the general forecaster captures the same conditional heteroscedasticity without committing to a parametric volatility law.

| Baseline                      | LL (all / cont.)   | CRPS (all / cont.) | Mean LL (cont.) |
|-------------------------------|--------------------|--------------------|-----------------|
| <b>AutoARIMA</b>              | <b>78 / 64 %</b>   | 51 / 47 %          | 2.09            |
| <b>AutoETS</b>                | <b>92 / 88 %</b>   | 68 / 67 %          | 2.20            |
| <b>SARIMAX</b> (statsmodels)  | <b>82 / 72 %</b>   | 53 / 49 %          | 2.13            |
| ETS (statsmodels)             | <b>93 / 89 %</b>   | 69 / 68 %          | 2.12            |
| <b>GARCH</b> (1, 1)- <i>t</i> | <b>79 / 67 %</b>   | 76 / 66 %          | <b>2.72</b>     |
| NF-Student <i>t</i>           | <b>100 / 100 %</b> | 78 / 76 %          | 0.82            |
| <b>AutoARIMA</b> + conformal  | <b>84 / 75 %</b>   | 37 / 29 %          | 1.47            |
| <b>AutoARIMA</b> + ACI        | <b>87 / 80 %</b>   | 36 / 28 %          | 1.89            |

Table 1: Per-series win-rate of **laplace** versus each baseline on 500 FRED change-series, all and on the 309 continuous series. LL counts series where **laplace** has the higher held-out log-likelihood; CRPS where it has the lower CRPS. The final column is each baseline’s mean continuous log-likelihood (**laplace**: **2.855**). Higher is better for LL; lower for CRPS.

**The neural result is not a scalp.** **laplace** beats NF-Student*t* on 100% of series, but we are deliberately loud about what this does and does not mean. NF here is **NeuralForecast** run in the *matched online univariate one-step* protocol — a tiny per-series multilayer perceptron with periodic refit — which is its **weak** regime. Neural and foundation forecasters [Salinas et al., 2020, Olivares et al., 2022] are built for *global cross-series* training, where they amortise structure across thousands of related series. Our result says nothing about **DeepAR**-style models in that home setting, and we do not claim it does; it says only that a single short univariate stream does not give a neural density head enough to beat an online ensemble. (We verified the NF densities are well formed — median log-likelihood 0.50, none floored — so the 100% is a real calibration gap, not a degenerate score.)

**CRPS.** As ever, CRPS tells a mixed story. **laplace** beats the smoothing baselines (**AutoETS**, ETS), **GARCH**-*t* and NF on CRPS, roughly ties the **ARIMA** family (**AutoARIMA** 51%, **SARIMAX** 53%), and *loses* to the conformal variants (37%, 36%), which are CRPS-optimised by construction. This is exactly the thesis of Sections 4: likelihood is the metric where a faithful density wins; CRPS is conformal’s home turf, where we are competitive but not dominant. A practitioner who truly only cares about CRPS can repoint the leaf (Section 9); a practitioner who cares about the density — for risk, sizing, or combination — should prefer the honest log-likelihood column.

**Scope.** Five hundred change-series, one-step horizon. Prophet, multi-step horizons, raw levels, and zero-shot foundation models [Ansari et al., 2024, Das et al., 2024, Woo et al., 2024, Rasul et al., 2023] — which require a fundamentally different, no-refit evaluation protocol — are left to an extended study.

## 9 Ablations: repointing the objective

A companion sweep on 2,501 FRED series isolates the effect of the terminal-leaf objective against the **crepes** conformal predictive system (given its three best calibration windows and scored on its own CDF via the pinball decomposition of CRPS). Because the conformal opponent there uses only a naive mean, we read this as an ablation of our own objective

knob rather than a state-of-the-art claim. Reporting per-policy and de-correlating to one vote per correlated family (195 families):

| Forecaster (identical structure)                                  | CRPS, raw | CRPS, family  | Mean LL     |
|-------------------------------------------------------------------|-----------|---------------|-------------|
| likelihood trunk + likelihood leaf                                | 82.6 %    | 48.7 %        | 2.96        |
| likelihood trunk + CRPS leaf ( <i>model first, conform last</i> ) | 89.5 %    | <b>60.1 %</b> | <b>3.01</b> |
| bare CRPS leaf (no trunk)                                         | 80.8 %    | 60.3 %        | 2.96        |

Table 2: Effect of the terminal-leaf objective. A general likelihood policy is essentially even with conformal on CRPS over independent families (48.7%); repointing only the terminal leaf to CRPS lifts that to 60.1% *and* raises log-likelihood, matching the dedicated CRPS specialist while modelling honestly.

Three further observations motivate the radically small public API. **Priors wash out:** across 25 series with  $\geq 2000$  changes the conventional prior-flavoured policies (trend, long-memory, drift, minimax, fast/slow) sit within 0.0006 nats of the uniform-prior general forecaster on the second half of each series, and the first-half winner beats the general forecaster on the second half only 44% of the time — worse than chance. **Conform-last is near-free:** on idealised light-tailed processes the CRPS leaf costs 0.005–0.007 nats versus the likelihood leaf, while on heavy-tailed and stochastic-volatility processes it *helps* (up to +0.04); real data lives in the second regime. **The lattice projection is free when unused.** Together these justify exposing exactly *one* general forecaster (uniform prior, CRPS leaf, lattice projection and coordinate learning on by default) plus *one* committed martingale specialist; the remaining structural ideas do not pay their way out of sample.

## 10 Discussion

The unifying theme is No-Free-Lunch [Wolpert and Macready, 1997] taken seriously. Every “method” is a prior; averaged over all series none dominates. The correct response is not to choose one but to make the priors explicit, weight them online, and bound the cost of a wrong prior by how fast the weighting abandons it. Read that way, the named methods of classical forecasting are convenience bundles, and the honest object is a single adaptive ensemble with two orthogonal terminal knobs — *which proper scoring rule the leaf optimises*, and *whether a lattice projection is active* — neither of which is a prior over the trunk. The empirical headline supports the design: against the strongest distributional opponents available on CPU, including the heavy-tail and neural state of the art, the general forecaster wins the likelihood race, and where it does not lead — CRPS — the deficit is confined to methods explicitly optimised for that metric and is recoverable by repointing a single knob.

## 11 Heritage and reproducibility

**skaters** is a from-scratch rewrite distilling ideas from **timemachines** [Cotton, 2021], an earlier competition-tested online-forecasting package, and builds on years of running live distributional-prediction contests at Microprediction, where forecasts are scored continuously as full distributions. The library is implemented twice — pure Python and zero-dependency JavaScript — and a parity suite checks roughly 90,000 probe values to  $10^{-6}$  on every run, so results are reproducible across both runtimes and in the browser via **Pyodide** [The Pyodide development team, 2021]. The benchmark of Section 8 is a

single parallel script (`benchmarks/sota_study.py`) that scores every method through the identical `Dist` interface; given a FRED key and the optional baselines (`statsforecast`, `arch`, `statsmodels`, `neuralforecast`) it reproduces Table 1 end to end.

## References

- Abdul Fatir Ansari et al. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024.
- Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- Tim Bollerslev. A conditionally heteroskedastic time series model for speculative prices and rates of return. *The Review of Economics and Statistics*, 69(3):542–547, 1987.
- George EP Box and David R Cox. An analysis of transformations. *Journal of the Royal Statistical Society: Series B*, 26(2):211–243, 1964.
- Peter Cotton. timemachines: Evaluation of online time-series forecasting algorithms. <https://github.com/microprediction/timemachines>, 2021.
- KE Dambis. On the decomposition of continuous submartingales. *Theory of Probability & Its Applications*, 10(3):401–410, 1965.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. *International Conference on Machine Learning*, 2024.
- Joseph L Doob. *Stochastic Processes*. Wiley, 1953.
- Lester E Dubins and Gideon Schwarz. On continuous martingales. *Proceedings of the National Academy of Sciences*, 53(5):913–916, 1965.
- Robert F Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of united kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- Federal Reserve Bank of St. Louis. Federal reserve economic data (FRED). <https://fred.stlouisfed.org>.
- Federico Garza, Max Mergenthaler Canseco, Cristian Challu, and Kin G Olivares. Stats-Forecast: Lightning fast forecasting with statistical and econometric models. *PyCon Salt Lake City, Utah, US 2022*, 2022. Nixtla.
- Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. *Advances in Neural Information Processing Systems*, 34:1660–1672, 2021.
- Tilmann Gneiting and Adrian E Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Jennifer A Hoeting, David Madigan, Adrian E Raftery, and Chris T Volinsky. Bayesian model averaging: a tutorial. *Statistical Science*, 14(4):382–401, 1999.
- Rob J Hyndman and Yeasmin Khandakar. Automatic time series forecasting: the forecast package for R. *Journal of Statistical Software*, 27(3):1–22, 2008.

- Rob J Hyndman, Anne B Koehler, Ralph D Snyder, and Simone Grose. A state space framework for automatic forecasting using exponential smoothing methods. *International Journal of Forecasting*, 18(3):439–454, 2002.
- John L Kelly. A new interpretation of information rate. *The Bell System Technical Journal*, 35(4):917–926, 1956.
- James E Matheson and Robert L Winkler. Scoring rules for continuous probability distributions. *Management Science*, 22(10):1087–1096, 1976.
- Kin G Olivares, Cristian Challu, Federico Garza, Max Mergenthaler Canseco, and Artur Dubrawski. NeuralForecast: User friendly state-of-the-art neural forecasting models. *PyCon Salt Lake City, Utah, US 2022*, 2022. Nixtla.
- Adrian E Raftery, Tilmann Gneiting, Fadoua Balabdaoui, and Michael Polakowski. Using Bayesian model averaging to calibrate forecast ensembles. *Monthly Weather Review*, 133(5):1155–1174, 2005.
- Kashif Rasul et al. Lag-Llama: Towards foundation models for probabilistic time series forecasting. *arXiv preprint arXiv:2310.08278*, 2023.
- David Salinas, Valentin Flunkert, Jan Gasthaus, and Tim Januschowski. DeepAR: Probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting*, 36(3):1181–1191, 2020.
- The Pyodide development team. Pyodide: Python with the scientific stack, compiled to WebAssembly. <https://pyodide.org>, 2021.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- Vladimir Vovk, Ilia Nouretdinov, Valery Manokhin, and Alexander Gammerman. Conformal predictive distributions with kernels. In *Conformal and Probabilistic Prediction and Applications*, 2019.
- BP Welford. Note on a method for calculating corrected sums of squares and products. *Technometrics*, 4(3):419–420, 1962.
- David H Wolpert and William G Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. *International Conference on Machine Learning*, 2024.
- In-Kwon Yeo and Richard A Johnson. A new family of power transformations to improve normality or symmetry. *Biometrika*, 87(4):954–959, 2000.