

A Stronger Univariate Benchmark for Foundational Time-Series Model Evaluation

Peter Cotton¹

<https://skaters.microprediction.org>

September 23, 2026

Abstract. A foundation model forecasts from a fixed window, with weights frozen at pretraining. An online forecaster carries state and revises it at every observation. Benchmarks for the first kind contain no example of the second: of the 132 entries on the GIFT-Eval leaderboard, five are not neural, and all five refit from scratch on each window. We propose `laplace`, a zero-dependency online distributional forecaster with no pretraining and no tunable parameters, as the missing reference. Scored one step ahead on held-out log density over an 18,815-series FRED panel, it beats TimesFM 3.0 on the median series in every stratum by roughly half a nat, and a post-hoc portfolio that sees only the model’s own issued density recovers most of the gap without retraining. The reference costs a `pip install` to include, and a Rust core with a bit-exact parity suite makes its numbers recompute on other hardware.

Keywords: time series, foundation models, benchmarking, probabilistic forecasting, proper scoring rules, online learning.

1 Introduction

A foundation model forecasts from a fixed window of recent history, with weights frozen at pretraining. An online forecaster carries state, and revises it with every observation it sees. Benchmarks for the first kind contain no example of the second.

Table 1 enumerates the classic arms of eleven studies, read from the papers rather than from their leaderboards: GIFT-Eval [Aksu et al., 2024], Chronos [Ansari et al., 2024], Chronos-2 [Ansari et al., 2025], TimesFM [Das et al., 2024], TiRex [Auer et al., 2025], Moirai [Woo et al., 2024], Moirai-MoE [Liu et al., 2024], fev-bench [Shchur et al., 2025], Lag-Llama [Rasul et al., 2023], TimeGPT-1 [Garza and Mergenthaler-Canseco, 2023], and the Monash archive [Godaheva et al., 2021a], which supplies several of the others’ baseline tables. Eighteen distinct methods appear across the eleven, and nine of those appear in exactly one paper. The working set is much smaller.

Three methods dominate. Automated ETS, ARIMA and Theta each appear in ten of the eleven. With seasonal naive they form a quartet that is the complete classic roster of six studies. No method appears in all eleven, though Theta does if its dynamically optimized variant is treated as the same model, which Garza and Mergenthaler-Canseco [2023] declines to do by listing both.

Five of the eleven ran no classical model at all for at least one headline comparison. Ansari et al. [2025] takes its classic numbers from Shchur et al. [2025] and the Aksu et al. [2024] leaderboard; Auer et al. [2025] takes its from the same leaderboard and from Ansari et al. [2024]; and Das et al. [2024], Woo et al. [2024] and Liu et al. [2024] inherit a Godaheva et al. [2021a] table, the last at two removes.

That matters because of what a Monash number is. Those models were run once, at library defaults, and the archive says so. Its ETS is forced non-seasonal above period 24 because the implementation cannot take longer cycles, and its harmonic regression is

Method	GIFT-Eval	Chronos	Chronos-2	TimesFM	TiReX	Moirai	Moirai-MoE	fev-bench	Lag-Llama	TimeGPT-1	Monash	n
(Auto)ETS	•	•	◦	◦	◦	◦	•	•	•	•	•	10
(Auto)ARIMA	•	•	◦	◦	◦	⊙	◦	•	•	•	•	10
(Auto)Theta	•	•	◦	◦	◦	◦	◦	•	•	•	•	10
Seasonal naive	•	•	◦	◦	•	•	•	•	•	•	•	8
Naive	•	•	◦	•	•	•	•	•	•	•	•	6
SES				◦	◦	◦					•	4
TBATS				◦	◦						•	3
DHR-ARIMA				◦	◦						•	3
SCUM ensemble	•	◦					•					3
DynOptTheta								•	•			2
Croston								•	•			2
AutoCES									•			1
ADIDA									•			1
IMAPA									•			1
Drift							•					1
HistoricAverage									•			1
ZeroModel									•			1
<i>NPTS</i>								•				1
laplace												0

Table 1: Every classic arm in eleven studies. • the authors ran it; ◦ the number is inherited from another paper’s table or from a leaderboard; ⊙ both, in different experiments. n counts the studies carrying each method. The two rows below the rule are the methods that carry state and issue a density. Tables inherited from [Godaheewa et al. \[2021a\]](#) print ARIMA and its harmonic-regression variant as one column applied to disjoint frequency groups; both rows are marked.

pinned to a single Fourier term. The table also reports point error only, so an inherited ETS column is a 2021 point-MAE number sitting beside a modern probabilistic score.

The arms that are run fare little better. [Ansari et al. \[2024\]](#) used library defaults with no tuning. [Aksu et al. \[2024\]](#) truncates statistical arms to 1000 observations and replaces any that exceeds a one-day limit with seasonal naive. [Shchur et al. \[2025\]](#) truncates to the same 1000 and caps the seasonal period at 200, and reports that its ensemble and automated ARIMA each failed four of a hundred tasks.

Statefulness and shape are separate matters, and neither is present. No arm in any of the eleven carries state between forecasts: the *statsforecast* predict path is handed an array, fits, forecasts, and keeps nothing. Seasonal naive, meanwhile, is a point forecaster serving as the probabilistic normalizer in seven of the eleven, and as the substituted score when another arm times out in two of them.

The last two rows of Table 1 are the methods in this literature that carry state and issue a density. *NPTS* is training-free and is a sampling distribution by construction, and it ships in both *GluonTS* and *AutoGluon*, which these studies import for their datasets, their metrics and their deep baselines. It is carried by one of the eleven [[Rasul et al., 2023](#)].

We scored *NPTS* through the same protocol as every other arm here. It loses to **laplace** on the median series in all four strata, by 0.05 nats on daily prices and 0.39 on

daily economics, over 1,500 series per stratum where the corpus allows. So the arm the field already ships and does not run is beatable, which is the weaker of the two things one might have found, and it sets the floor this paper’s reference has to clear.

This paper supplies such an arm and reports what it does to the comparison. The contribution is deliberately small: a reference, a protocol, and four tables.

2 `laplace`: an online distributional forecaster

`laplace` imports nothing, requires no GPU, performs no fitting step, and exposes no tunable hyperparameters. One parameter-free call returns a predictive density.

The predictive object is a weighted mixture of Gaussians. That choice buys three things a reference arm needs: the mixture is closed under affine maps and under mixing, it admits a closed-form CRPS, and it approximates heavy tails as a scale mixture. A density can therefore be pushed through an invertible map and scored against a proper rule [Gneiting and Raftery, 2007] without sampling.

A *transform* is a stateful causal map that consumes one observation and emits one transformed value, paired with an inverse that rebuilds predictive distributions for future observations from the transformed coordinate, given the state it has accumulated. Differencing, standardisation, online autoregression, *GARCH*-style scaling, seasonal differencing and the Yeo–Johnson coordinates are all transforms. *Conjugation* composes one with an inner forecaster: the observation is transformed, the inner model issues a density in the new coordinate, and the inverse carries that density back. Multi-step uncertainty accumulates through the inverse rather than by simulation.

`laplace` is one fixed composition of these parts, with no knob exposed. Candidate models enter it only as mean forecasters, weighted online by predictive likelihood, so that a shaped leaf governs the tail rather than the dispersion of candidate means. A lattice projection places atoms on the exact values a grid-quoted series revisits, and vanishes identically on continuous data. Beyond one step the predictive mixes over decimated clocks, since the fine clock wins short horizons and a coarse one wins long. A conditional tail fit splices the extreme quantiles, which a bulk fit leaves too thin.

Cotton [2026] derives all of this. We restate only what a reader needs to judge the comparison.

Two properties matter for a reference arm, and neither is accuracy. The first is cost. Adding `laplace` to a leaderboard is a `pip install` and one function call, with no weights to download, no device to allocate, and no context length to choose. The pure-Python implementation, 280 KB of source across 28 files, is the reference; a Rust core runs the same arithmetic about twenty times faster.

Table 2 gives the runtime, and the two regimes disagree. Cold is one forecast for a series not seen before, which is what a benchmark harness asks for. Warm is one new observation on a series already tracked, which is what a deployment asks for.

Warm is the regime that separates the arms. One online update costs 0.012 milliseconds, against 7.44 for the cheapest foundation model here and 329.83 for the largest.

That per-forecast figure compounds, because a benchmark asks for a forecast at every step of every series. Table 3 multiplies it out over the 1,204,288 one-step forecasts this panel contains.

The spread runs from fifteen seconds to four and a half days. It also accounts for the shape of the suites in Table 1: scoring a stateless model at every step of a large panel is expensive enough that one forecast origin per series is the affordable design, and four of the eleven use exactly that.

Model	cold (ms)		warm (ms)
	batch 1	batch 64	
<code>laplace</code> (Rust)	3.39	—	0.012
<code>laplace</code> (Python)	67.08	—	0.251
AutoETS (refit)	1.93	—	1.93
Prophet (refit)	29.97	—	29.97
AutoARIMA (refit)	62.13	—	62.13
Chronos-Bolt-small	7.44	0.72	7.44
TiRex	96.47	76.34	95.68
TimesFM-2.5-200M	323.44	5.10	329.83

Table 2: Milliseconds per forecast, context 256, horizon 12, 64 series, one thread, Apple silicon, *skaters* 0.16.0. Foundation models batch in the cold regime, so the batched column is reported alongside the per-series one. In the warm regime `laplace` performs one state update; a stateless model re-runs its forward pass over a shifted window and a refitting baseline re-estimates from scratch, so for those the warm cost equals the cold cost.

Model	ms/step	batch 1	batch 64
<code>laplace</code> (Rust)	0.012	15 s	—
<code>laplace</code> (Python)	0.251	5 min	—
AutoETS (refit)	1.93	39 min	—
Chronos-Bolt-small	7.44	2.5 h	14 min
Prophet (refit)	29.97	10.0 h	—
AutoARIMA (refit)	62.13	20.8 h	—
TiRex	95.68	32.0 h	25.5 h
TimesFM-2.5-200M	329.83	4.6 d	1.7 h

Table 3: Wall clock to score the whole panel one step at a time, single threaded. A stateless model re-runs its forward pass at every step, so its per-step cost is its cold cost, and the batch-64 column is the fair figure for an offline sweep that can run many series in lockstep.

The second is reproducibility across machines. The Rust core uses portable mathematics (*libm*) and produces identical bits on x86 and ARM, and a shared parity suite checks roughly 100,000 probe values to 10^{-6} per implementation. A number one person publishes therefore recomputes on another person’s hardware, which is the property a denominator needs and a pretrained checkpoint behind an inference stack does not have.

A denominator sets what a skill score means, and a weak one flatters everything divided by it.

Sections 4 and 5 measure how strong it is. `laplace` takes the per-series log-likelihood against every classical arm in Table 1, and against TimesFM 3.0 in every stratum tested. Cotton [2026] reports the same ordering against neural and foundation-model baselines across the non-price FRED universe; *GARCH-t* on asset prices is the one arm that escapes, and that split is reported rather than averaged away.

3 Protocol

Every arm sees the same series, the same windows and the same scoring rule. The context is 128 observations, the test window rolls one step ahead, and no arm is fitted to the evaluation data. The score is held-out log predictive density, differenced per series so that

Arm	LL win		CRPS win		N
	raw	family	raw	family	
AutoARIMA (<i>statsforecast</i>)	87 %	87 %	58 %	63 %	1688
AutoETS (<i>statsforecast</i>)	95 %	94 %	77 %	78 %	1688
Theta (R)	92 %	95 %	74 %	79 %	1917
auto.arima (R)	84 %	85 %	51 %	60 %	1917

Table 4: Share of series on which **laplace** scores better, over 5,397 continuous non-price FRED series in 1,213 correlated families. The family figure averages within a family first, so that a cluster of related series counts once. N is the subset each arm covers under its compute budget. Every arm forecasts one step ahead at every point, conditioned on all data up to that point, with coefficients re-estimated every 25 steps.

each series’ intrinsic difficulty cancels.

We report the median of that difference and the per-series win, draw and loss counts from a Diebold–Mariano test with HAC standard errors and a draw band, following Diebold and Mariano [1995]. We do not report a pooled mean. Averaging a log density across heterogeneous series estimates nothing, because a few near-constant series dominate the sum; Koning et al. [2005] made the corresponding argument about the M3 competition, and Hewamalage et al. [2023] prescribe per-series testing.

The context is fixed rather than growing. Several foundation models degrade beyond the window they were pretrained on, and most of them silently truncate over-long input rather than raising, so feeding four thousand points to a model configured for two thousand scores a two-thousand-context model without saying so. A protocol that lets context grow needs a declared per-model cap, set at the documented architectural maximum, and recorded per arm.

Recording it matters for a second reason. The Woo et al. [2024] benchmark numbers come from a per-dataset search over context lengths from one thousand to five thousand, jointly with patch size, which the arms it is compared against do not receive.

4 Against the classic arms

The arms of Table 1 are the denominator every published skill score divides by, so their strength sets what a skill score means. Table 4 scores **laplace** against them on the non-price FRED universe.

On log-likelihood the margin is wide: the closest arm loses on 84 % of series and the widest on 95 %. On CRPS the arms do better, which is the expected pattern, since CRPS is forgiving of a mis-shaped tail that a log score punishes.

These arms are stronger than the ones the studies of Table 1 run.² Their harnesses re-estimate at each window, but the forecast for step j of a window conditions only on data up to the window origin, so inside a window of length h the baseline never sees the observations it is forecasting past. Four of the eleven use a single origin per series.

²The two differ in which half of the model goes stale. *statsforecast*’s **refit** argument freezes the fitted coefficients between refits but hands **forward** the current history, so the arms here are coefficient-stale and conditioning-current. The frozen parameters include the innovation variance, so the interval inherits a scale up to 25 steps old, which is the one setting that could flatter Table 4. Re-estimating at every step instead, on twenty paired series, moves AutoARIMA by a median of 0.005 nats and AutoETS by -0.000 , against margins of 0.31 to 0.59 nats. The published harnesses are the mirror image: they re-estimate at every window and freeze the conditioning.

Stratum	N	Win / draw / loss	Median ΔLL	Loss rate
Daily, economic	2,179	66/923/1,190	−0.854	54.6 %
Daily, price/returns	7,633	15/4,593/3,025	−0.616	39.6 %
Weekly, economic	3,013	517/1,481/1,015	−0.520	33.7 %
Monthly, economic	5,576	193/3,592/1,791	−0.647	32.1 %
M4-hourly, seasonal	414	8/241/165	−0.543	39.9 %

Table 5: TimesFM 3.0 against `laplace`, one step ahead on held-out log density, 128-length context. ΔLL is the per-series difference, model minus reference, reported as a median. The loss rate is the share of series on which the model loses.

Stratum	N	Win / draw / loss	Median ΔLL	Loss rate
Daily, economic	2,179	76/1,383/720	−0.094	33.0 %
Daily, price/returns	7,633	152/7,136/345	−0.008	4.5 %
Weekly, economic	3,011	467/2,045/499	−0.049	16.6 %
Monthly, economic	5,195	350/4,480/365	−0.056	7.0 %
M4-hourly, seasonal	414	58/352/4	−0.007	1.0 %

Table 6: The same model wrapped in a never-worse portfolio with the reference. The loss rate falls by as much as thirty-nine points; the median stays negative in every stratum.

This is the sense in which the incumbent denominators are soft. A skill score computed against them credits a model for clearing a bar that a dependency-free online method clears by a wide margin on the same data.

5 Against TimesFM 3.0

Table 5 scores TimesFM 3.0 [Das et al., 2024] zero-shot against the reference on an 18,815-series FRED panel. Negative favours the reference.

The model loses on the majority of series in every stratum, by a median of roughly half a nat. Half a nat is a large gap in log-likelihood terms.

Table 6 repeats the comparison with the model wrapped in a post-hoc portfolio that combines its own issued predictive with the reference. The wrap requires no retraining and no access to model internals; it sees only the density the model published.

The wrapped model is safe to deploy, since its loss rate collapses. It remains behind the reference, since the median stays negative in every stratum.

A post-hoc combination that sees only the published density recovers most of the gap. At this horizon the pretraining is therefore contributing less than the raw comparison suggests, and what it does contribute survives a wrapper.

6 The reconstruction moves the answer more than the model does

A quantile-only model does not emit a density, and a density must exist before a log score does. Something must build one from the quantiles, and that something is a modelling choice independent of the model being scored.

Fix TimesFM 3.0 and fix 3,200 held-out test points, then vary only the reconstruction across four methods. The mean log density moves from −0.373 to +2.104. The mean per-point spread is 3.13 nats, which exceeds the entire measured gap between the model

and the reference in every stratum of Table 5. Two methods already in independent use elsewhere differ by 0.62 nats.

The choice therefore belongs in the protocol, stated rather than assumed. It is not peculiar to this study: the Aksu et al. [2024] submission interface requires a mean and nine deciles, so a model emitting a genuine density must reduce it before scoring and a model emitting samples must convert.

That interface’s interval score asks for the 2.5th and 97.5th percentiles, and a sample forecast returns a nearest order statistic without interpolation. At the twenty samples the reference *Chronos* notebook uses, those two quantiles are the minimum and the maximum of twenty draws.

7 Horizon, grid, and statefulness

The results above do not contradict any published leaderboard, because no published leaderboard occupies this regime. Three settings separate them.

The horizon never reaches one step. The shortest prediction length in Aksu et al. [2024] is six, and Shchur et al. [2025] use domain horizons such as thirty and one hundred sixty-eight.

The scoring rule is a nine-point grid. What these suites report as CRPS is a mean of weighted quantile losses at the deciles, so nothing outside the tenth and ninetieth percentile enters any published score. Ansari et al. [2025] call that grid a defect and move to twenty-one levels including the first and ninety-ninth.

And no baseline carries state, which is the hole of Section 1. Each setting is defensible for the task those suites set themselves. Together they define a regime in which an online distributional method cannot compete.

8 Scope

Everything measured here is univariate and one step ahead. A single pretrained model amortized across many related series can exploit cross-series and panel structure that no univariate method reaches by construction, and that is where pretraining should pay for itself most directly. This paper does not make that comparison, and it is the more decisive one.

Leakage is controlled in this literature by dataset name rather than by time. Ansari et al. [2024] state the temporal standard, that evaluation data should begin after the last pretraining observation, and then decline it as a minimal risk. Meyer et al. [2025] and Oreshkin et al. [2026] report that a corpus used in TimesFM pretraining reappears in its Monash evaluation; their pretraining table lists Traffic at 862 series and 15,122,928 observations, and the Monash archive file [Godaheewa et al., 2021b] holds 862 series of exactly 17,544 steps.

A calibration framing would have to engage Adler et al. [2026], who find these models better calibrated than automated ARIMA and N-BEATS on standard datasets. Nothing here contradicts that. Calibration on a decile grid and sharpness of a full density are different measurements.

Finally, the panel is FRED, and FRED is the universe these models’ pretraining corpora almost entirely omit: the corpus of Woo et al. [2024] is 0.09 % economic by observation, and the entire economics domain of Aksu et al. [2024] is the M4 collection. Whether the gap survives on the energy and transport data these models were built for is an open question, and Wang et al. [2026] suggest it will not.

9 Reproduction

The paper directory holds four files. `constants.py` derives every number in Sections 5 and 6 from a version-controlled store. `verify_paper.py` re-derives each one, compares it against the tables above, and exits non-zero on drift. `REPRODUCE.md` gives three tiers: the figures check in seconds from committed files, the summaries rebuild from per-step records in minutes, and the records regenerate from raw data in days.

The data vintage is published as a release asset in the repository `skaters-benchmarking-data`. FRED revises series, so a fresh pull returns different vintages and will not reproduce these numbers. The archive carries 210,574 public-domain series; commercial index series are excluded and listed, which costs the price stratum and nothing else.

References

- Coen Adler, Tiffany Chang, Fabian Draxler, Samar Abdi, and Padhraic Smyth. Beyond accuracy: Are time series foundation models well-calibrated? *International Conference on Learning Representations*, 2026. arXiv:2510.16060.
- Taha Aksu, Gerald Woo, Juncheng Liu, Xu Liu, Chenghao Liu, Silvio Savarese, Caiming Xiong, and Doyen Sahoo. GIFT-Eval: A benchmark for general time series forecasting model evaluation. *arXiv preprint arXiv:2410.10393*, 2024.
- Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, et al. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024. arXiv:2403.07815.
- Abdul Fatir Ansari, Oleksandr Shchur, Caner Turkmen, et al. Chronos-2: From univariate to universal forecasting. *arXiv preprint arXiv:2510.15821*, 2025.
- Andreas Auer et al. TiRex: Zero-shot forecasting across long and short horizons with enhanced in-context learning. *arXiv preprint arXiv:2505.23719*, 2025.
- Peter Cotton. Transforms all the way down: Automatic online distributional forecasting by conjugation. Technical report, SSRN, 2026. <https://skaters.microprediction.org>.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *PMLR*, pages 10148–10167, 2024. arXiv:2310.10688.
- Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business and Economic Statistics*, 13(3):253–263, 1995.
- Azul Garza and Max Mergenthaler-Canseco. TimeGPT-1. *arXiv preprint arXiv:2310.03589*, 2023.
- Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Rakshitha Godahewa, Christoph Bergmeir, Geoffrey I. Webb, Rob J. Hyndman, and Pablo Montero-Manso. Monash time series forecasting archive. *NeurIPS Datasets and Benchmarks Track*, 2021a. arXiv:2105.06643.

- Rakshitha Godahewa, Christoph Bergmeir, Geoffrey I. Webb, Rob J. Hyndman, and Pablo Montero-Manso. Traffic hourly dataset. Zenodo record 4656132, 2021b.
- Hansika Hewamalage, Klaus Ackermann, and Christoph Bergmeir. Forecast evaluation for data scientists: Common pitfalls and best practices. *Data Mining and Knowledge Discovery*, 37:788–832, 2023.
- Alex J. Koning, Philip Hans Franses, Michèle Hibon, and H. O. Stekler. The M3 competition: Statistical tests of the results. *International Journal of Forecasting*, 21(3): 397–409, 2005.
- Xu Liu, Juncheng Liu, Gerald Woo, et al. Moirai-MoE: Empowering time series foundation models with sparse mixture of experts. *arXiv preprint arXiv:2410.10469*, 2024.
- Lars Meyer, David Kaltenpoth, Jonas Zalipski, and Emmanuel Müller. Rethinking evaluation in the era of time series foundation models: (un)known information leakage challenges. *arXiv preprint arXiv:2510.13654*, 2025.
- Boris N. Oreshkin, Aakash Jauhari, Rohan Selvam, et al. Zero-shot forecasting by simulation alone. *arXiv preprint arXiv:2601.00970*, 2026.
- Kashif Rasul, Arjun Ashok, Andrew Robert Williams, et al. Lag-llama: Towards foundation models for probabilistic time series forecasting. *arXiv preprint arXiv:2310.08278*, 2023.
- Oleksandr Shchur, Abdul Fatir Ansari, Caner Turkmen, Lorenzo Stella, Nick Erickson, et al. fev-bench: A realistic benchmark for time series forecasting. *arXiv preprint arXiv:2509.26468*, 2025.
- Y. Wang, X. Quan, Z. Yang, and A. Srivastava. Spectral predictability and the limits of zero-shot time series forecasting. *arXiv preprint arXiv:2511.08884*, 2026.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *Proceedings of the 41st International Conference on Machine Learning*, 2024. arXiv:2402.02592.